

Dimensions of Responsible AI

Generative AI

Module 2 – Lesson 3

Today's activities



- Review of foundations of responsible AI
- Dimensions of responsible AI
- Privacy and Security
- Robustness
- Veracity
- Fairness and safety
- Transparency and explainability
- Governance
- Controllability

Review: Core dimensions of responsible AI

Controllability

Having mechanisms to monitor and steer AI system behavior

Privacy & Security

Appropriately obtaining, using and protecting data and models

Safety

Preventing harmful system output and misuse

Fairness

Considering impacts on different groups of stakeholders

Veracity & Robustness

Achieving correct system outputs, even with unexpected or adversarial inputs

Explainability

Understanding and evaluating outputs generated by an AI system

Transparency

Enabling stakeholders to make informed choices about their engagement with an AI system

Governance

Incorporating best practices into the AI supply chain, including providers and deployers

Privacy and Security

Privacy and security

Risk:

- The possibility of an **adverse event with consequences** on any or multiple core dimensions of Responsible AI

Security

Exposure to threats that can **compromise the integrity, confidentiality, or availability** of an ML/AI system

Privacy


Exposure or mishandling of **sensitive or personal data** in the context of an interaction with an ML/AI system

Privacy in the context of generative AI

- Foundation models have been trained with unprecedented amounts of data
- Model users might not know exactly which type of data went into training a pre-trained model

Security in the context of generative AI

- Input to generative AI models is mediated via **prompts**
- Thus, security incidents in generative AI are related to **unsafe interactions between user prompt, the model, and its output**



Robustness

Robustness


- It is the model's ability to:
 - **Maintain performance** and **stability** under various challenging conditions
 - **Generalize well** and **perform reliably** in real-world data deviating from training data

Robustness

It is the model's ability to generalize well and perform reliably in real-world data deviating from training data


Types of Robustness: Noise

- Model can handle noisy or corrupted input data. [\[Jia et al. \(2017\)\]](#)
- Example:




Types of Robustness: OOD

- **OOD**: Out-of-distribution/ Out-of-domain
- Models find it harder to maintain performance on topics that are scarce in the training data [\[Hendryks et al. \(2021\)\]](#)
- Example: LLM pre-trained on social media data. LLM scarcely trains on math problems. (long-tail)



Types of robustness: Adversarial

- Model can withstand adversarial attacks that try to mislead the model. [[Goodfellow et al. \(2015\)](#)]
- Example: LLM is able to respond appropriately even when the input has been **deliberately and intentionally modified to mislead the model**.



Types of robustness: Multi-task

- Model can perform well on a diverse set of tasks. [[McCann et al. \(2018\)](#)]
- Example: A single LLM which can perform multiple different natural language processing tasks **without significant performance degradation** such as:
 - Question answering
 - Sentiment analysis
 - Language translation
- Benchmark Dataset: Massive Multitask Language Understanding (MMLU) benchmark can be found [here](#).

Veracity

What is veracity?

- Improving veracity means **reducing** hallucinations
- **Hallucinations**: Plausible-sounding but **factually incorrect** information due to the probabilistic nature of AI outputs

Veracity


Truthfulness and accuracy of the generated responses from LLMs

Why is veracity important?

- LLMs are pre-trained on vast amounts of data that can contain **inaccuracies, biases**, and sometimes **even intentionally misleading** information
- High veracity improves the LLM response to be more **reliable, factual, and trustworthy** for tasks such as:
 - Question-answering
 - Content generation
 - Decision support systems


Types of veracity (1/3)

1. **Credibility:** Evaluates the trustworthiness and reliability of the source providing the information. Compare answers against multiple sources.
2. **Factuality:** Measures the extent to which a claim or statement is supported by empirical evidence and objective facts.




Types of veracity (2/3)

3. Coherence : Different parts of model response **should not contradict** each other.



4. Completeness : Model response is **comprehensive**: It covers all relevant aspects of the topic



Types of veracity (3/3)

5. Temporal : Model response is sequentially accurate (**chronological**) and consistent regarding timing of events.

Input prompt:
"What were the key events leading up to the end of World War II in Europe?"



LLM response:
"The key events leading up to the end of World War II in Europe include D-Day in June 1944, the Battle of the Bulge in **December 1944**, the liberation of Paris in **August 1944**, and the fall of Berlin in May 1945."

Fairness and Safety

Fairness

- Fairness considers how a system **impacts different subpopulations** of users.
- Fairness can generally be described as the mitigation of unintended bias.
- Bias (in the responsible AI space) means disparity in performance.

Fairness

Considering impacts on different groups of stakeholders

Safety

- Preventing undesirable outputs
 - Such as toxic, hurtful and personal statements
- Misusing the system for unintended purposes:
 - Jailbreaking, prompt injection, adversarial attacks etc

Safety

Preventing harmful system output and misuse.

Transparency and Explainability

Transparency

- **Providing details** about the AI system
 - Technical reports
 - Model cards/details
- Make the user aware when they are interacting with an AI system

Transparency

Enabling stakeholders to make informed choices
about their engagement with an AI system

Explainability

- The LLM can **provide explanations** for its predictions, helping users understand the model's **decision-making process** and build trust in the model's outputs. [\[Jacovi et al. \(2020\)\]](#)
- Example: In Retrieval Augmented Generation (RAG) workflows, LLM can provide **citations/source attributions** of the exact text used to generate the final LLM response.

Explainability

Understanding and evaluating
outputs generated by an AI system

Example: Citations and source attributions

- Here we see **citations/source attributions** of the exact text used to generate the final LLM response in Amazon Knowledge Bases when using the LLM to generate the final response.

The screenshot displays the 'Test Knowledge base' interface. At the top, there's a toggle for 'Generate responses' and the model name 'Claude v2.1 | ODT'. Below this, a user query is entered: 'what is the northwind schema like?'. The response from the LLM is: 'The Northwind database schema contains several tables like customers, products, orders, order details, employees, shippers, suppliers, territories, etc. It models a typical ecommerce database with en etc. [1]'. A citation box is shown next to the response, containing the link '[1] northwind_ddl.sql' and a close button. A tooltip is visible over the citation, showing a list of examples of SQL queries that can answer questions from the Northwind database, such as 'Get subtotal for each order' and 'select OrderID, format(sum(UnitPrice...'. A link '< Hide result details' is also present.

Test Knowledge base

☒ Generate responses

Claude v2.1 | ODT

[Change model](#)

what is the northwind schema like ?

The Northwind database schema contains several tables like customers, products, orders, order details, employees, shippers, suppliers, territories, etc. It models a typical ecommerce database with en etc. [1]

[\[1\] northwind_ddl.sql](#)

< Hide result details

-- <EXAMPLES> -- Few Examples of SQLite sql queries that can answer questions from Northwind database are: -- Question: Get subtotal for each order. -- Query select OrderID, format(sum(UnitPrice...

Governance

What is governance?

- Provides **oversight and control** to manage risks and achieve objectives
- Supports **ethical, transparent, and responsible** decision-making
- Promotes **accountability** and builds stakeholder trust
- Enables **compliance** with laws, regulations, and industry standards

Governance

Governance refers to the systems, processes, and structures by which an organization is directed, controlled, and held accountable

Controllability

Controllability

- **Monitor** and **tune** the AI system to maintain desired performance

Controllability

Having mechanisms
to monitor and steer AI system behavior

Steering the AI system's behavior


- **In-context solutions**
 - Provide details, examples and guidance via prompts
- **Fine-tuning**
 - Update the model's weights to adjust the outputs of the model
 - Model alignment
- **Compound systems**
 - Combining various elements together to improve quality and performance of the system
 - Multiple calls to LLMs, retrievers, tools, agents, etc.


Monitor AI System

- Evaluate models and applications
 - **Covered in Lesson 1 of this module: Evaluating LLMs**
- AI systems need continuous monitoring
 - Evaluate at regular frequencies
 - Performance may decline with drifts in data, policies, project scope, etc.

Next lesson

- This lesson covered various dimensions of responsible AI
- In the next lesson, you will learn about techniques to improve security and safety





Thank you!